

Lab6

57118233 周厚霖

Task 1: Implementing a Simple Firewall

Task 1.A: Implement a Simple Kernel Module

原始目录存在空格，目录的空格被 make 识别为编译的 target，所以我们需要把 kernel_module 拷贝到 /home/seed/ 目录下进行编译。

```
[07/22/21]seed@VM:~/kernel_module$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M]  /home/seed/kernel_module/hello.o
  Building modules, stage 2.
  MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/kernel_module/hello.o
see include/linux/module.h for more information
  CC [M]  /home/seed/kernel_module/hello.mod.o
  LD [M]  /home/seed/kernel_module/hello.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

编译成功后测试以下命令。

```
[07/22/21]seed@VM:~/kernel_module$ sudo rmmod hello.ko
[07/22/21]seed@VM:~/kernel_module$ sudo insmod hello.ko
[07/22/21]seed@VM:~/kernel_module$ lsmod | grep hello
hello                16384  0
[07/22/21]seed@VM:~/kernel_module$ sudo rmmod hello.ko
```

```
[07/22/21]seed@VM:~/kernel_module$ dmesg | grep World
[ 4615.481235] Hello World!
[ 4672.813707] Bye-bye World!
[ 4674.396915] Hello World!
[ 4690.904285] Bye-bye World!
```

Task 1.B: Implement a Simple Firewall Using Netfilter

1、和前面一样，将文件拷贝到 /home/seed/ 下面进行编译。

```
[07/25/21]seed@VM:~/packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M]  /home/seed/packet_filter/seedFilter.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC [M]  /home/seed/packet_filter/seedFilter.mod.o
  LD [M]  /home/seed/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

加载内核前，可以看到 dig @8.8.8.8 www.example 命令可以得到响应。

```
[07/25/21]seed@VM:~/packet_filter$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 61984
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                20013   IN      A      93.184.216.34

;; Query time: 48 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Sun Jul 25 05:01:12 EDT 2021
;; MSG SIZE rcvd: 60
```

加载到内核后，可以看到防火墙生效。

```
[07/25/21]seed@VM:~/packet_filter$ sudo insmod seedFilter.ko
[07/25/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
seedFilter                16384  0
[07/25/21]seed@VM:~/packet_filter$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; connection timed out; no servers could be reached
```

最后从内核中移除。

```
[07/25/21]seed@VM:~/packet_filter$ sudo rmmod seedFilter
[07/25/21]seed@VM:~/packet_filter$ lsmod | grep seedFilter
```

2、在进行实验时，每次修改完代码，需要重新 make 编译，然后使用 `sudo insmod seedFilter.ko` 加载到内核，可以使用 `lsmod | grep seedFilter` 查看模块是否在内核，进行 `dig @8.8.8.8 www.example.com` 操作后，可使用 `sudo dmesg -c` 查看信息，每次测试后，需要运行 `sudo rmmod seedFilter` 从内核中移除模块。

数据报从进入系统，进行 IP 校验以后，首先经过第一个 HOOK 函数 `NF_INET_PRE_ROUTING` 进行处理，然后就进入路由代码，其决定该数据报是需要转发还是发给本机的。


```
[07/25/21]seed@VM:~/packet_filter$ sudo dmesg -c
[18536.191183] e1000: ens33 NIC Link is Down
[18540.223729] e1000: ens33 NIC Link is Up 1000 Mbps Full Duplex, Flow Control: None
[18984.105489] e1000: ens33 NIC Link is Down
[18990.144009] e1000: ens33 NIC Link is Up 1000 Mbps Full Duplex, Flow Control: None
[19040.278364] Registering filters.
[19052.284669] *** PRE_ROUTING
[19052.284671] 192.168.23.2 --> 192.168.23.130 (UDP)
[19052.285798] *** PRE_ROUTING
[19052.285800] 127.0.0.1 --> 127.0.0.53 (UDP)
[19052.288420] *** PRE_ROUTING
[19052.288421] 192.168.23.2 --> 192.168.23.130 (UDP)
[19052.288558] *** PRE_ROUTING
[19052.288559] 127.0.0.53 --> 127.0.0.1 (UDP)
[19052.288623] *** PRE_ROUTING
[19052.288624] 127.0.0.1 --> 127.0.0.53 (UDP)
[19052.290981] *** PRE_ROUTING
[19052.291013] 192.168.23.2 --> 192.168.23.130 (UDP)
[19052.293504] *** PRE_ROUTING
[19052.293506] 192.168.23.2 --> 192.168.23.130 (UDP)
[19052.293769] *** PRE_ROUTING
[19052.293770] 127.0.0.53 --> 127.0.0.1 (UDP)
```

```
[19064.293893] *** PRE_ROUTING
[19064.293896] 127.0.0.1 --> 127.0.0.1 (UDP)
[19064.294076] *** Dropping 8.8.8.8 (UDP), port 53
[19069.295665] *** Dropping 8.8.8.8 (UDP), port 53
[19074.299678] *** Dropping 8.8.8.8 (UDP), port 53
[19116.318325] *** PRE_ROUTING
[19116.318328] 192.168.23.2 --> 192.168.23.130 (UDP)
[19116.318980] *** PRE_ROUTING
[19116.318982] 127.0.0.1 --> 127.0.0.53 (UDP)
[19116.321402] *** PRE_ROUTING
[19116.321403] 192.168.23.2 --> 192.168.23.130 (UDP)
[19116.321766] *** PRE_ROUTING
[19116.321767] 127.0.0.53 --> 127.0.0.1 (UDP)
[19116.321830] *** PRE_ROUTING
[19116.321831] 127.0.0.1 --> 127.0.0.53 (UDP)
[19116.324237] *** PRE_ROUTING
[19116.324239] 192.168.23.2 --> 192.168.23.130 (UDP)
[19116.326414] *** PRE_ROUTING
[19116.326416] 192.168.23.2 --> 192.168.23.130 (UDP)
[19116.326539] *** PRE_ROUTING
[19116.326540] 127.0.0.53 --> 127.0.0.1 (UDP)
[19117.342518] *** PRE_ROUTING
[19117.342525] 192.168.23.130 --> 224.0.0.251 (UDP)
```

若该数据报是发被本机的，则该数据经过 `HOOK` 函数 `NF_INET_LOCAL_IN` 处理以后然后传递给上层协议。

```
[07/25/21]seed@VM:~/packet_filter$ sudo dmesg -c
[19341.048679] The filters are being removed.
[19346.140769] Registering filters.
[19348.832838] *** LOCAL_IN
[19348.832840] 127.0.0.1 --> 127.0.0.1 (UDP)
[19348.833055] *** Dropping 8.8.8.8 (UDP), port 53
[19353.832834] *** Dropping 8.8.8.8 (UDP), port 53
[19358.836846] *** Dropping 8.8.8.8 (UDP), port 53
```

若该数据报应该被转发则它被 `NF_INET_FORWARD` 处理。

```
[07/25/21]seed@VM:~/packet_filter$ sudo dmesg -c
[19502.900875] The filters are being removed.
[19604.819118] Registering filters.
[19618.667010] *** Dropping 8.8.8.8 (UDP), port 53
[19623.669846] *** Dropping 8.8.8.8 (UDP), port 53
[19628.673838] *** Dropping 8.8.8.8 (UDP), port 53
```



```

struct iphdr *iph;
struct udphdr *udph;

u16 port = 53;
char ip[16] = "8.8.8.8";
u32 ip_addr;

if (!skb) return NF_ACCEPT;

iph = ip_hdr(skb);
// Convert the IPv4 address from dotted decimal to 32-bit binary
in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);

if (iph->protocol == IPPROTO_UDP) {
    udph = udp_hdr(skb);
    if (iph->daddr == ip_addr && ntohs(udph->dest) == port){
        printk(KERN_WARNING "**** Dropping %pI4 (UDP), port %d\n", &(iph->daddr), port);
        return NF_DROP;
    }
}
return NF_ACCEPT;
}

unsigned int blockTCP(void *priv, struct sk_buff *skb,
                      const struct nf_hook_state *state)
{
    struct iphdr *iph;
    struct tcphdr *tcph;

    u16 port = 23;
    char ip[16] = "10.9.0.1";
    u32 ip_addr;

    if (!skb) return NF_ACCEPT;

    iph = ip_hdr(skb);
    // Convert the IPv4 address from dotted decimal to 32-bit binary
    in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);

    if (iph->protocol == IPPROTO_TCP) {
        tcph = tcp_hdr(skb);
        if (iph->daddr == ip_addr && ntohs(tcph->dest) == port){
            printk(KERN_WARNING "**** Dropping %pI4 (TCP), port %d\n", &(iph->daddr), port);
            return NF_DROP;
        }
    }
    return NF_ACCEPT;
}

unsigned int blockICMP(void *priv, struct sk_buff *skb,
                      const struct nf_hook_state *state)
{
    struct iphdr *iph;
    struct icmphdr *icmph;

    char ip[16] = "10.9.0.1";

```

```

u32 ip_addr;

if (!skb) return NF_ACCEPT;

iph = ip_hdr(skb);
// Convert the IPv4 address from dotted decimal to 32-bit binary
in4_pton(ip, -1, (u8 *)&ip_addr, '\0', NULL);

if (iph->protocol == IPPROTO_ICMP) {
    icmp_h = icmp_hdr(skb);
    if (iph->daddr == ip_addr){
        printk(KERN_WARNING "*** Dropping %pI4 (ICMP)\n", &(iph->daddr));
        return NF_DROP;
    }
}
return NF_ACCEPT;
}

unsigned int printInfo(void *priv, struct sk_buff *skb,
                      const struct nf_hook_state *state)
{
    struct iphdr *iph;
    char *hook;
    char *protocol;

    switch (state->hook){
        case NF_INET_LOCAL_IN:      hook = "LOCAL_IN";      break;
        case NF_INET_LOCAL_OUT:     hook = "LOCAL_OUT";     break;
        case NF_INET_PRE_ROUTING:   hook = "PRE_ROUTING";   break;
        case NF_INET_POST_ROUTING: hook = "POST_ROUTING";   break;
        case NF_INET_FORWARD:       hook = "FORWARD";       break;
        default:                    hook = "IMPOSSIBLE";     break;
    }
    printk(KERN_INFO "*** %s\n", hook); // Print out the hook info

    iph = ip_hdr(skb);
    switch (iph->protocol){
        case IPPROTO_UDP: protocol = "UDP"; break;
        case IPPROTO_TCP: protocol = "TCP"; break;
        case IPPROTO_ICMP: protocol = "ICMP"; break;
        default:           protocol = "OTHER"; break;
    }
    // Print out the IP addresses and protocol
    printk(KERN_INFO "    %pI4 --> %pI4 (%s)\n",
              &(iph->saddr), &(iph->daddr), protocol);

    return NF_ACCEPT;
}

int registerFilter(void) {
    printk(KERN_INFO "Registering filters.\n");

    hook1.hook = printInfo;
    hook1.hooknum = NF_INET_LOCAL_OUT;
}

```

```

hook1.pf = PF_INET;
hook1.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook1);

hook2.hook = blockUDP;
hook2.hooknum = NF_INET_POST_ROUTING;
hook2.pf = PF_INET;
hook2.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook2);

hook3.hook = blockICMP;
hook3.hooknum = NF_INET_PRE_ROUTING;
hook3.pf = PF_INET;
hook3.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook3);

hook4.hook = blockTCP;
hook4.hooknum = NF_INET_PRE_ROUTING;
hook4.pf = PF_INET;
hook4.priority = NF_IP_PRI_FIRST;
nf_register_net_hook(&init_net, &hook4);

return 0;
}

void removeFilter(void) {
    printk(KERN_INFO "The filters are being removed.\n");
    nf_unregister_net_hook(&init_net, &hook1);
    nf_unregister_net_hook(&init_net, &hook2);
    nf_unregister_net_hook(&init_net, &hook3);
    nf_unregister_net_hook(&init_net, &hook4);
}

module_init(registerFilter);
module_exit(removeFilter);

MODULE_LICENSE("GPL");

```

开启容器，登录到 docker1(10.9.0.5)，在容器上分别进行 ping 10.9.0.1 和 telnet 10.9.0.1。

```

root@886cc22ed637:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^C
--- 10.9.0.1 ping statistics ---
5 packets transmitted, 0 received, 100% packet loss, time 4088ms

root@886cc22ed637:/# telnet 10.9.0.1
Trying 10.9.0.1...
^C

```

在本机上查看内核缓存。

```

[07/25/21]seed@VM:~/packet_filter$ sudo dmesg -c
[23142.330556] Registering filters.
[23160.504894] *** Dropping 10.9.0.1 (ICMP)
[23161.522619] *** Dropping 10.9.0.1 (ICMP)
[23162.545505] *** Dropping 10.9.0.1 (ICMP)
[23163.569967] *** Dropping 10.9.0.1 (ICMP)
[23164.593234] *** Dropping 10.9.0.1 (ICMP)
[23166.718789] *** Dropping 10.9.0.1 (TCP), port 23
[23167.730388] *** Dropping 10.9.0.1 (TCP), port 23
[23169.746973] *** Dropping 10.9.0.1 (TCP), port 23

```


Task 2: Experimenting with Stateless Firewall Rules

每个任务前都要清理 `table` 或重启路由器的 `docker`。

Task 2.A: Protecting the Router

输入以下命令后, `ping 10.9.0.11` 和 `telnet 10.9.0.11` 都不通。

```
root@b16b9b75497b:/# iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@b16b9b75497b:/# iptables -A OUTPUT -p icmp --icmp-type echo-request -j ACCEPT
root@b16b9b75497b:/# iptables -P OUTPUT DROP
root@b16b9b75497b:/# iptables -P INPUT DROP
```

```
root@99c053c27260:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
^C
--- 10.9.0.11 ping statistics ---
10 packets transmitted, 0 received, 100% packet loss, time 9214ms

root@99c053c27260:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
```

修改成这样, 可以 `ping` 通, 但是 `telnet` 不通。(题目中给出的应该是搞反了)

```
root@b16b9b75497b:/# iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@b16b9b75497b:/# iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
root@b16b9b75497b:/# iptables -P OUTPUT DROP
root@b16b9b75497b:/# iptables -P INPUT DROP
```

允许其他主机 `ping` 通防火墙

```
root@b16b9b75497b:/# iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@b16b9b75497b:/# iptables -A INPUT -p icmp --icmp-type echo-request -j
ACCEPT
```

设置 `INPUT` 和 `OUTPUT` 链默认为丢包

```
root@b16b9b75497b:/# iptables -P OUTPUT DROP
root@b16b9b75497b:/# iptables -P INPUT DROP
```

```
root@99c053c27260:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.391 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.138 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.176 ms
^C
--- 10.9.0.11 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2054ms
rtt min/avg/max/mdev = 0.138/0.235/0.391/0.111 ms
root@99c053c27260:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
```

Task 2.B: Protecting the Internal Network


```
# 内部主机可以ping通外部主机
root@b16b9b75497b:/# iptables -A FORWARD -p icmp --icmp-type echo-request -d 10.9.0.5/24 -j ACCEPT
root@b16b9b75497b:/# iptables -A FORWARD -p icmp --icmp-type echo-reply -d 192.168.60.0/24 -j ACCEPT
# 外部主机不能ping通内部主机
root@b16b9b75497b:/# iptables -A FORWARD -p icmp --icmp-type echo-request -d 192.168.60.0/24 -j DROP
# 路由器接受ping命令
root@b16b9b75497b:/# iptables -A INPUT -p icmp -j ACCEPT
root@b16b9b75497b:/# iptables -A OUTPUT -p icmp -j ACCEPT
# 修改策略为丢弃所以数据包
root@b16b9b75497b:/# iptables -P FORWARD DROP
```

设置如下:

```
root@b16b9b75497b:/# iptables -L
Chain INPUT (policy ACCEPT)
target prot opt source destination
ACCEPT icmp -- anywhere anywhere

Chain FORWARD (policy DROP)
target prot opt source destination
ACCEPT icmp -- anywhere 10.9.0.0/24 icmp echo-request
ACCEPT icmp -- anywhere 192.168.60.0/24 icmp echo-reply
DROP icmp -- anywhere 192.168.60.0/24 icmp echo-request

Chain OUTPUT (policy ACCEPT)
target prot opt source destination
ACCEPT icmp -- anywhere anywhere
```

从外部主机ping路由器, 可以ping通; ping内部主机, 不通; telnet内部主机, 不通。

```
root@99c053c27260:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.157 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.101 ms
^C
--- 10.9.0.11 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1021ms
rtt min/avg/max/mdev = 0.101/0.129/0.157/0.028 ms
root@99c053c27260:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
^C
--- 192.168.60.5 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2048ms

root@99c053c27260:/# telnet 192.168.60.5
Trying 192.168.60.5...
^C
```

内部主机ping外部主机, 可以ping通; telnet外部主机, 不通。

```
root@8c08b332f530:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=0.158 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.132 ms
^C
--- 10.9.0.5 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1015ms
rtt min/avg/max/mdev = 0.132/0.145/0.158/0.013 ms
root@8c08b332f530:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
```

Task 2.C: Protecting Internal Servers

```
root@b16b9b75497b:/# iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 -j ACCEPT
root@b16b9b75497b:/# iptables -A FORWARD -p tcp --sport 23 -s 192.168.60.5 -j ACCEPT
root@b16b9b75497b:/# iptables -A FORWARD -d 10.9.0.0/24 -j DROP
root@b16b9b75497b:/# iptables -A FORWARD -d 192.168.60.0/24 -j DROP
```

设置如下:

```
root@b16b9b75497b:/# iptables -L
Chain INPUT (policy ACCEPT)
target     prot opt source               destination

Chain FORWARD (policy ACCEPT)
target     prot opt source               destination
ACCEPT     tcp  --  anywhere             host1-192.168.60.5.net-192.168.60.0  tcp dpt:telnet
ACCEPT     tcp  --  host1-192.168.60.5.net-192.168.60.0  anywhere             tcp spt:telnet
DROP       all  --  anywhere             10.9.0.0/24
DROP       all  --  anywhere             192.168.60.0/24

Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination
```

从外部主机 (10.9.0.5) telnet 192.168.60.5, 可以连接成功。

```
root@99c053c27260:/# telnet 192.168.60.5
Trying 192.168.60.5...
^C
root@99c053c27260:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
8c08b332f530 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage
```

从外部主机 (10.9.0.5) telnet 192.168.60.6, 无法连接。

```
root@99c053c27260:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
```

从内部主机 (192.168.60.5) telnet 10.9.0.5, 无法连接, 内部主机 (192.168.60.5) telnet 192.168.60.6, 连接成功。

```
root@8c08b332f530:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
root@8c08b332f530:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
fe5bfdbac545 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage
```

Task 3: Connection Tracking and Stateful Firewall

Task 3.A: Experiment with the Connection Tracking

ICMP 的连接状态保持时间只有 30 秒左右。

```
root@b16b9b75497b:/# conntrack -L
icmp      1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
icmp      1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
icmp      1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
icmp      1 24 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
icmp      1 18 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
icmp      1 14 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
icmp      1 8 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=46 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=46 mark=
0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
```

UDP 的连接状态保持时间和也只有 20~30 秒之间。

```
root@b16b9b75497b:/# conntrack -L
udp       17 24 src=10.9.0.5 dst=192.168.60.5 sport=55350 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=909
0 dport=55350 mark=0 use=1
udp       17 0 src=10.9.0.5 dst=192.168.60.5 sport=33888 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090
 dport=33888 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 2 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
udp       17 16 src=10.9.0.5 dst=192.168.60.5 sport=55350 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=909
0 dport=55350 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
udp       17 9 src=10.9.0.5 dst=192.168.60.5 sport=55350 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090
 dport=55350 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@b16b9b75497b:/# conntrack -L
udp       17 3 src=10.9.0.5 dst=192.168.60.5 sport=55350 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090
 dport=55350 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

TCP 的连接状态保持时间非常长，大约 430000 秒。

```
root@c36907490561:/# conntrack -L
tcp       6 431989 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=38630 dport=90
90 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=38630 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@c36907490561:/# conntrack -L
tcp       6 431986 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=38630 dport=90
90 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=38630 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@c36907490561:/#
```

Task 3.B: Setting Up a Stateful Firewall

```
root@c36907490561:/# iptables -F
root@c36907490561:/# iptables -A FORWARD -p tcp -m conntrack --ctstate
ESTABLISHED,RELATED -j ACCEPT
root@c36907490561:/# iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 --syn
-m conntrack --ctstate NEW -j ACCEPT
root@c36907490561:/# iptables -A FORWARD -p tcp --dport 23 -d 10.9.0.0/24 --syn
-m conntrack --ctstate NEW -j ACCEPT
root@c36907490561:/# iptables -P FORWARD DROP
```


从外部主机 (10.9.0.5) telnet 192.168.60.5 和 192.168.0.6 , 均可以连接成功。

```
root@7ae2f1716d9e:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
7d88e3ad0d23 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

```
root@7ae2f1716d9e:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
```

从内部主机 (192.168.60.5) telnet 10.9.0.5 和 192.168.60.6 , 连接成功。

```
root@7d88e3ad0d23:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
7ae2f1716d9e login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

```
root@7d88e3ad0d23:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
0603b13e4bea login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

Task 4: Limiting Network Traffic

```
root@c36907490561:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minut -
-limit-burst 5 -j ACCEPT
root@c36907490561:/# iptables -A FORWARD -s 10.9.0.5 -j DROP
```

可以观察到前六个包的速度很快, 后面每隔6秒发一个包。

```
root@7ae2f1716d9e:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.152 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.096 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.220 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.127 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.219 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.225 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.223 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.229 ms
^C
--- 192.168.60.5 ping statistics ---
20 packets transmitted, 8 received, 60% packet loss, time 19464ms
rtt min/avg/max/mdev = 0.096/0.186/0.229/0.049 ms
```

如果只执行第一条命令，从外部 (10.9.0.5) ping 192.168.60.5，可以观察到和平时的发包速度一样，因为 iptables 默认的 FORWARD 表是接受所有包，所以如果不写第二条命令，发包会正常进行。

Task 5: Load Balancing

```
root@c36907490561:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode nth --every 3 --packet 0 -j DNAT --to-destination
192.168.60.5:8080
root@c36907490561:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode nth --every 3 --packet 1 -j DNAT --to-destination
192.168.60.6:8080
root@c36907490561:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode nth --every 3 --packet 2 -j DNAT --to-destination
192.168.60.7:8080
```

按顺序 hello_1 被发送到 192.168.60.5 8080，hello_2 被发送到 192.168.60.6 8080，hello_3 被发送到 192.168.60.7 8080。

```
root@a4cff2dbff4a:/# nc -luk 8080
hello_1
Labs_20.04
```

```
root@7d88e3ad0d23:/# nc -luk 8080
hello_2
```

```
root@0603b13e4bea:/# nc -luk 8080
hello_3
Labs_20.04
```

```
root@c36907490561:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode random --probability 0.33 -j DNAT --to-destination
192.168.60.5:8080
root@c36907490561:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode random --probability 0.33 -j DNAT --to-destination
192.168.60.6:8080
root@c36907490561:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m
statistic --mode random --probability 0.34 -j DNAT --to-destination
192.168.60.7:8080
```

虽然是等概率发送数据，但每个主机收到的数量各不相同，甚至有的差异较大，当样本数量足够多时，应该是趋于平均的。

```
root@7d88e3ad0d23:/# nc -luk 8080
hello_7
hello_12
hello_13
hello_15
```

```
root@0603b13e4bea:/# nc -luk 8080
hello_1
hello_2
hello_5
hello_6
hello_9
hello_10
hello_11
hello_14
```

```
root@a4cff2dbff4a:/# nc -luk 8080
hello_3
hello_4
hello_8
```